

Apify Integration PRD & Task List

Project Context

Product: LinkedIn Ads Intelligence Platform — SaaS tool for agencies to research competitor LinkedIn ads.

Tech stack: Next.js 14, Supabase (PostgreSQL), Clerk (auth), Prisma, Tailwind CSS v3, shadcn/ui, Trigger.dev

Current state: Project initialized, Prisma schema exists with Advertiser and Ad models, Clerk auth configured, database tables pushed to Supabase.

We are integrating the **Apify LinkedIn Ad Library Scraper** (Actor ID: `AdwCDVyhFcWXQF9tg`, slug: `silva95gustavo/linkedin-ad-library-scraper`) to scrape ads from LinkedIn's Ad Library. This replaces the originally planned custom Playwright scraper.

1. Apify Actor Details

Input Parameters

The actor accepts the following input:

```
json

{
  "startUrls": [
    "https://www.linkedin.com/ad-library/search?companyId=2027242"
  ],
  "resultsLimit": null,
  "skipDetails": false,
  "proxy": {
    "useApifyProxy": true,
    "apifyProxyGroups": ["RESIDENTIAL"],
    "apifyProxyCountry": "US"
  }
}
```

- `startUrls` (required): Array of LinkedIn Ad Library search URLs. Each URL targets one advertiser using `companyId=<LINKEDIN_COMPANY_ID>`.
- `resultsLimit` (optional): Max number of ads to return. Leave null/empty for all.
- `skipDetails` (optional, boolean): If true, skips fetching full ad details (less data, cheaper). We want `false`.
- Proxy: Datacenter-Automatic with US country works fine.

Run Configuration

- Timeout: 3600 seconds (default)
- Memory: 512 MB (default)

- Cost: \$2.19 for 547 ads from one advertiser (\$0.004/ad)
- At MVP scale (5-10 advertisers, ~2000 ads/month): estimated \$8-20/month

API Endpoint

To start a run via API:

```
POST https://api.apify.com/v2/acts/silva95gustavo~linkedin-ad-library-scraper/runs
```

To get dataset results after run completes:

```
GET https://api.apify.com/v2/datasets/{datasetId}/items?format=json
```

Authentication: Pass `token` query parameter with your Apify API token.

2. Apify Output Schema (Real Data)

Below is the complete field inventory based on a real scrape of 547 ads. Fields are categorized by availability across ad formats.

Universal Fields (present in ALL 9 formats)

Apify Field	Type	Example	Notes
<code>adId</code>	string	"1145280734"	Unique ad identifier from LinkedIn
<code>adLibraryUrl</code>	string	"https://www.linkedin.com/ad-library/detail/1145280734"	Direct link to ad on LinkedIn
<code>advertiserName</code>	string	"Simplicate"	Company name
<code>advertiserUrl</code>	string	"https://www.linkedin.com/company/2027242"	LinkedIn company page URL
<code>advertiserLogo</code>	string	"https://media.licdn.com/dms/image/..."	Logo image URL
<code>format</code>	string	"SINGLE_IMAGE"	One of 9 format values (see below)
<code>availability.start</code>	string (date)	"2026-02-12"	Ad start date (YYYY-MM-DD)
<code>availability.end</code>	string (date)	"2026-02-17"	Ad end date (YYYY-MM-DD)

Apify Field	Type	Example	Notes
impressions	string	"1k-5k"	Impression range as text
paidBy	string	"Simplicate"	Entity that paid — sometimes differs from advertiser (e.g. agency)
ctas	string[]	["Meer informatie"]	Call-to-action buttons. Can be empty array.
targeting	object	{"language": "Nederlands", "location": "Netherlands"}	Has language (optional) and location (optional)
startUrl	string	"https://www.linkedin.com/ad-library/search?companyIds=2027242"	The search URL used to find this ad

Common Fields (present in most but not all formats)

Apify Field	Type	Present In	Notes
body	string	All except FOLLOW_COMPANY	Ad copy text. Can be very long with emojis/unicode.
clickUrl	string	All except MESSAGE, DOCUMENT	Destination URL when ad is clicked.
headline	string	SINGLE_IMAGE, VIDEO, DOCUMENT, SPOTLIGHT	Short headline text.
impressionsPerCountry	array	SINGLE_IMAGE, CAROUSEL, VIDEO, MESSAGE, DOCUMENT, EVENT, TEXT	Array of {country: string, impressions: string}

Format-Specific Fields

Apify Field	Type	Format	Example
imageUrl	string	SINGLE_IMAGE	Single image URL
videoUrl	string	VIDEO	Video file URL (.mp4)
documentUrl	string	DOCUMENT	PDF document URL
imageUrls	string[]	DOCUMENT	Array of page preview image URLs
slides	array	CAROUSEL	Array of {imageUrl: string} objects

Format Values Found in Real Data

Format Value	Count (547 ads)	Has Media	Media Field(s)
SINGLE_IMAGE	475 (87%)	Yes	imageUrl
CAROUSEL	21	Yes	slides[].imageUrl
VIDEO	15	Yes	videoUrl
MESSAGE	13	No	— (InMail ad, no creative shown)
TEXT	10	No	— (text-only ad)
DOCUMENT	6	Yes	documentUrl, imageUrls[]
EVENT	5	No	— (links to LinkedIn event)
SPOTLIGHT	2	No	— (dynamic ad, personalized per viewer)
FOLLOW_COMPANY	1	No	— (dynamic follower ad)

Not seen in this dataset but may exist: LEAD_GEN_FORM, CONVERSATION. Handle these as unknown/other.

Impression Values (string ranges)

The impressions field is always a string. These are the possible values observed:

```
"< 1k", "<\u200a1k", "1k-5k", "5k-10k", "10k-20k",
"20k-30k", "30k-50k", "50k-100k", "100k-150k", "150k-200k"
```

Note: Some values contain a Unicode thin space (\u200a) before "1k". Normalize these when storing.

3. Prisma Schema Updates

Update the existing Prisma schema. The Advertiser model stays mostly the same. The Ad model needs significant updates to match real Apify data.

Updated Advertiser Model

```
prisma
```

```
model Advertiser {
    id      String @id @default(cuid())
    name    String
    linkedinCompanyId String @unique // e.g. "2027242"
    linkedinUrl   String?          // e.g. "https://www.linkedin.com/company/2027242"
    logoUrl     String?          // from Apify advertiserLogo field
    status      String @default("approved") // "pending" | "approved" | "rejected"
    lastScrapedAt DateTime?
    totalAdsFound Int   @default(0)
    createdAt    DateTime @default(now())
    updatedAt    DateTime @updatedAt

    ads         Ad[]
    scrapeRuns  ScrapeRun[]
}
```

Updated Ad Model

```
prisma
```

```

model Ad {
    id          String  @id @default(cuid())
    externalId   String  @unique    // Apify "adId" — LinkedIn's ad ID
    advertiserId String
    advertiser   Advertiser @relation(fields: [advertiserId], references: [id], onDelete: Cascade)

    // Core content
    format       String      // "SINGLE_IMAGE" | "CAROUSEL" | "VIDEO" | etc.
    bodyText     String?     // Apify "body" — ad copy
    headline     String?     // Apify "headline"
    callToAction String?     // First element of Apify "ctas" array
    destinationUrl String?   // Apify "clickUrl"

    // Media — primary URL for quick access
    mediaUrl     String?     // Primary image/video URL (imageUrl, videoUrl, or first slide)
    mediaData     Json?       // Format-specific media data (slides array, imageUrls array, documentUrl, etc.)

    // Dates and performance
    startDate    DateTime?   // Apify "availability.start"
    endDate      DateTime?   // Apify "availability.end"
    impressions  String?     // Raw string like "1k-5k"

    // Targeting
    targetLanguage String?    // Apify "targeting.language"
    targetLocation String?    // Apify "targeting.location"

    // Metadata
    paidBy        String?     // Apify "paidBy" — may differ from advertiser
    adLibraryUrl String?     // Direct link to LinkedIn Ad Library detail page
    impressionsPerCountry Json? // Array of {country, impressions} objects

    // Timestamps
    firstSeenAt   DateTime @default(now()) // When we first scraped this ad
    lastSeenAt    DateTime @default(now()) // Updated each scrape where ad still appears
    createdAt     DateTime @default(now())
    updatedAt     DateTime @updatedAt

    @@index([advertiserId])
    @@index([format])
    @@index([startDate])
    @@index([endDate])
}

```

New ScrapeRun Model (for tracking/debugging)

```
model ScrapeRun {  
    id      String  @id @default(cuid())  
    advertiserId  String  
    advertiser  Advertiser @relation(fields: [advertiserId], references: [id], onDelete: Cascade)  
    apifyRunId  String?          // Apify's run ID for reference  
    apifyDatasetId String?        // Apify's dataset ID  
    status     String  @default("pending") // "pending" | "running" | "completed" | "failed"  
    adsFound   Int    @default(0)  
    adsNew     Int    @default(0)    // Ads not previously in our database  
    adsUpdated  Int    @default(0)    // Ads that already existed and were refreshed  
    costUsd    Float?         // Estimated cost of this run  
    errorMessage  String?  
    startedAt   DateTime @default(now())  
    completedAt  DateTime?  
  
    @@index([advertiserId])  
    @@index([status])  
}
```

4. Field Mapping: Apify → Prisma

This is the single transformation layer. All Apify field names should ONLY appear in the mapping function. The rest of the codebase only uses our Prisma field names.

typescript

```
// lib/apify/transform.ts

interface ApifyAd {
    adId: string;
    adLibraryUrl: string;
    advertiserName: string;
    advertiserUrl: string;
    advertiserLogo: string;
    format: string;
    body?: string;
    headline?: string;
    clickUrl?: string;
    ctas: string[];
    availability: { start: string; end: string };
    impressions: string;
    paidBy: string;
    targeting: { language?: string; location?: string };
    imageUrl?: string;
    videoUrl?: string;
    documentUrl?: string;
    imageUrls?: string[];
    slides?: Array<{ imageUrl: string }>;
    impressionsPerCountry?: Array<{ country: string; impressions: string }>;
    startUrl: string;
}

// Returns data ready for prisma.ad.upsert()
function transformApifyAd(raw: ApifyAd, advertiserId: string) {
    return {
        externalId: raw.adId,
        advertiserId,
        format: raw.format,
        bodyText: raw.body ?? null,
        headline: raw.headline ?? null,
        callToAction: raw.ctas?.[0] ?? null,
        destinationUrl: raw.clickUrl ?? null,
        mediaUrl: getMediaUrl(raw),
        mediaData: getMediaData(raw),
        startDate: raw.availability?.start ? new Date(raw.availability.start) : null,
        endDate: raw.availability?.end ? new Date(raw.availability.end) : null,
        impressions: normalizeImpressions(raw.impressions),
        targetLanguage: raw.targeting?.language ?? null,
        targetLocation: raw.targeting?.location ?? null,
        paidBy: raw.paidBy ?? null,
        adLibraryUrl: raw.adLibraryUrl ?? null,
        impressionsPerCountry: raw.impressionsPerCountry ?? null,
    }
}
```

```

    lastSeenAt: new Date(),
};

}

// Extract the primary media URL based on format
function getMediaUrl(raw: ApifyAd): string | null {
    switch (raw.format) {
        case 'SINGLE_IMAGE': return raw.imageUrl ?? null;
        case 'VIDEO': return raw.videoUrl ?? null;
        case 'CAROUSEL': return raw.slides?.[0]?.imageUrl ?? null;
        case 'DOCUMENT': return raw.imageUrls?.[0] ?? null;
        default: return null;
    }
}

// Store format-specific media as structured JSON
function getMediaData(raw: ApifyAd): object | null {
    switch (raw.format) {
        case 'CAROUSEL':
            return { slides: raw.slides ?? [] };
        case 'DOCUMENT':
            return { documentUrl: raw.documentUrl, imageUrls: raw.imageUrls ?? [] };
        case 'VIDEO':
            return { videoUrl: raw.videoUrl };
        case 'SINGLE_IMAGE':
            return { imageUrl: raw.imageUrl };
        default:
            return null;
    }
}

// Normalize Unicode thin spaces and whitespace in impression strings
function normalizeImpressions(raw: string): string {
    return raw?.replace(/\u200a\u200b\u00a0/g, "").trim() ?? null;
}

```

5. Apify Client Service

typescript

```
// lib/apify/client.ts

const APIFY_TOKEN = process.env.APIFY_API_TOKEN;
const ACTOR_ID = 'silva95gustavo~linkedin-ad-library-scraper';
const BASE_URL = 'https://api.apify.com/v2';

// Start a scrape run for one advertiser
async function startScrapeRun(linkedinCompanyId: string): Promise<{ runId: string }> {
    const response = await fetch(
        `${BASE_URL}/acts/${ACTOR_ID}/runs?token=${APIFY_TOKEN}`,
        {
            method: 'POST',
            headers: { 'Content-Type': 'application/json' },
            body: JSON.stringify({
                startUrls: [
                    `https://www.linkedin.com/ad-library/search?companyIds=${linkedinCompanyId}`
                ],
                skipDetails: false,
            }),
        }
    );
    const data = await response.json();
    return { runId: data.data.id };
}

// Poll until run finishes (or use webhook — see Task 4)
async function waitForRun(runId: string): Promise<{ datasetId: string; status: string }> {
    // Poll every 10 seconds, timeout after 20 minutes
    const maxAttempts = 120;
    for (let i = 0; i < maxAttempts; i++) {
        const response = await fetch(
            `${BASE_URL}/actor-runs/${runId}?token=${APIFY_TOKEN}`
        );
        const data = await response.json();
        const status = data.data.status;

        if (status === 'SUCCEEDED') {
            return { datasetId: data.data.defaultDatasetId, status };
        }
        if (status === 'FAILED' || status === 'ABORTED' || status === 'TIMED-OUT') {
            throw new Error(`Apify run ${runId} ended with status: ${status}`);
        }
        await new Promise(resolve => setTimeout(resolve, 10_000));
    }
    throw new Error(`Apify run ${runId} timed out waiting`);
}
```

```
// Fetch all results from a completed run's dataset
async function getDatasetItems(datasetId: string): Promise<ApifyAd[]> {
    const response = await fetch(
        `${BASE_URL}/datasets/${datasetId}/items?format=json&token=${APIFY_TOKEN}`
    );
    return response.json();
}
```

6. Upsert Strategy

When storing ads, use `upsert` on `externalId` to handle re-scraping the same advertiser:

typescript

```

async function storeAds(ads: ApifyAd[], advertiserId: string): Promise<{ new: number; updated: number }> {
    let newCount = 0;
    let updatedCount = 0;

    for (const raw of ads) {
        const data = transformApifyAd(raw, advertiserId);

        const existing = await prisma.ad.findUnique({
            where: { externalId: data.externalId },
            select: { id: true },
        });

        await prisma.ad.upsert({
            where: { externalId: data.externalId },
            create: {
                ...data,
                firstSeenAt: new Date(),
            },
            update: {
                ...data,
                // Don't overwrite firstSeenAt on update
            },
        });
    }

    if (existing) {
        updatedCount++;
    } else {
        newCount++;
    }
}

return { new: newCount, updated: updatedCount };
}

```

7. Environment Variables

Add to `.env.local`:

```
APIFY_API_TOKEN=your_apify_api_token_here
```

Add to `.env.example`:

```
APIFY_API_TOKEN=
```

8. Task List (Ordered)

Complete these tasks in order. Each task should be a working increment — test before moving to the next.

Task 1: Update Prisma Schema

Files to modify: `prisma/schema.prisma`

1. Update the `Advertiser` model to match the schema in Section 3. Add fields: `linkedinCompanyId` (unique), `linkedinUrl`, `logoUrl`, `status`, `lastScrapedAt`, `totalAdsFound`.
2. Replace the existing `Ad` model with the schema in Section 3. Key: `externalId` must be `@unique`.
3. Add the new `ScrapeRun` model from Section 3.
4. Run `npx prisma db push` to apply changes.
5. Run `npx prisma generate` to update the Prisma client.

Verify: Open Prisma Studio (`npx prisma studio`) and confirm all three tables exist with correct columns.

Task 2: Create Apify Type Definitions

Files to create: `lib/apify/types.ts`

1. Define the `ApifyAd` interface matching the raw JSON output (see Section 4).
2. Define an `ApifyRunStatus` type: `'READY' | 'RUNNING' | 'SUCCEEDED' | 'FAILED' | 'ABORTED' | 'TIMED-OUT'`.
3. Define an `AdFormat` type with all known values: `'SINGLE_IMAGE' | 'CAROUSEL' | 'VIDEO' | 'MESSAGE' | 'TEXT' | 'DOCUMENT' | 'EVENT' | 'SPOTLIGHT' | 'FOLLOW_COMPANY'`.
4. Export all types.

Task 3: Create Transform Layer

Files to create: `lib/apify/transform.ts`

1. Implement `transformApifyAd()` exactly as shown in Section 4.
2. Implement `getMediaUrl()` — extracts the single most relevant media URL per format.
3. Implement `getMediaData()` — stores all format-specific media as a JSON object.
4. Implement `normalizeImpressions()` — strips Unicode thin spaces.
5. Export all functions.

Important: This is the ONLY file that references Apify field names. Every other file works with our Prisma types.

Task 4: Create Apify Client

Files to create: `lib/apify/client.ts`

1. Implement `startScrapeRun(linkedinCompanyId)` as shown in Section 5.
2. Implement `waitForRun(runId)` with polling (10s interval, 20min timeout).
3. Implement `getDatasetItems(datasetId)` to fetch results.
4. Implement a convenience function `scrapeAdvertiser(linkedinCompanyId)` that chains: `startScrapeRun → waitForRun → getDatasetItems → return transformed ads`.
5. Add proper error handling: wrap all fetch calls in try/catch, check `response.ok`, throw descriptive errors.

Task 5: Create Ad Storage Service

Files to create: `lib/services/ad-storage.ts`

1. Implement `storeAds(apifyAds, advertiserId)` using the upsert strategy from Section 6.
2. After storing ads, update the `Advertiser` record: set `lastScrapedAt` to now, set `totalAdsFound` to total ad count.
3. Create a `ScrapeRun` record to log the run (status, counts, timing).
4. Return a summary: `{ adsNew, adsUpdated, totalProcessed }`.

Task 6: Create Scrape API Route (Manual Trigger)

Files to create: `app/api/scrape/[advertiserId]/route.ts`

1. Create a POST endpoint that accepts an `advertiserId`.
2. Look up the advertiser in the database, get their `linkedinCompanyId`.
3. Create a `ScrapeRun` record with status "running".
4. Call `scrapeAdvertiser()` from the Apify client.
5. Call `storeAds()` to persist results.
6. Update the `ScrapeRun` record with results (status "completed", ad counts).
7. Return JSON response with scrape results.
8. On error: update `ScrapeRun` to "failed" with error message, return 500.

Note: This is a long-running request (1-5 minutes). For MVP this is acceptable when triggered manually from admin UI. Later, this moves to Trigger.dev background jobs.

Task 7: Test End-to-End with Real Data

No new files — manual testing.

1. Add a test advertiser to the database: Simplicate, `linkedinCompanyId: 2027242`.
2. Trigger the scrape API route.

3. Verify in Prisma Studio:

- 547 ads were created in the Ad table
- All 9 format types are represented
- `mediaUrl` is populated for SINGLE_IMAGE, VIDEO, CAROUSEL, DOCUMENT formats
- `mediaData` JSON contains the correct format-specific data
- `startDate` and `endDate` are valid dates
- `impressions` has no Unicode artifacts
- The ScrapeRun record shows correct counts

4. Trigger the scrape again and verify upsert works: 0 new, 547 updated, `lastSeenAt` timestamps refreshed.

Task 8: Admin Page — Advertiser Management

Files to create: `app/admin/advertisers/page.tsx`, related components

1. List all advertisers with: name, status, lastScrapedAt, totalAdsFound.
2. "Add Advertiser" form with fields: name, linkedinCompanyId.
3. A "Scrape Now" button per advertiser that calls the API route from Task 6.
4. Show scrape status/progress (loading state while scrape runs).
5. Display last scrape result (new/updated counts) after completion.

Task 9: Budget Tracking (Cost Guard Rail)

Files to modify: `lib/services/ad-storage.ts`, ScrapeRun model

1. After each run completes, estimate cost: $(adsFound * 0.004)$ (based on ~\$0.004/ad from real data).
2. Store `costUsd` in the ScrapeRun record.
3. Add a helper `getMonthlySpend()` that sums ScrapeRun costs for the current month.
4. Before starting a new run, check if monthly spend exceeds a configurable limit (default: \$50). If exceeded, throw an error instead of starting the run.
5. Show monthly spend on the admin page.

9. Example Apify Output (Real Data)

SINGLE_IMAGE

json

```
{  
  "adId": "1145280734",  
  "adLibraryUrl": "https://www.linkedin.com/ad-library/detail/1145280734",  
  "advertiserLogo": "https://media.licdn.com/dms/image/...",  
  "body": "[VACATURE] Customer Support Consultant bij Simplicate...",  
  "clickUrl": "https://werkenbij.simplicate.nl/o/support-consultant-2?...",  
  "ctas": ["Meer informatie"],  
  "advertiserName": "Simplicate",  
  "advertiserUrl": "https://www.linkedin.com/company/2027242",  
  "availability": { "start": "2026-02-12", "end": "2026-02-17" },  
  "format": "SINGLE_IMAGE",  
  "paidBy": "Simplicate",  
  "impressions": "<\u200a1k",  
  "targeting": { "language": "Nederlands", "location": "Amsterdam Area, Groningen en omgeving" },  
  "headline": "Vacature Support Consultant | Amsterdam of Groningen",  
  "imageUrl": "https://media.licdn.com/dms/image/...",  
  "startUrl": "https://www.linkedin.com/ad-library/search?companyId=2027242"  
}
```

CAROUSEL

```
json  
{  
  "adId": "750982734",  
  "format": "CAROUSEL",  
  "body": "...",  
  "impressions": "< 1k",  
  "slides": [  
    { "imageUrl": "https://media.licdn.com/dms/image/.../ssu-carousel-card-image_lax_1200_1200/..." },  
    { "imageUrl": "https://media.licdn.com/dms/image/.../ssu-carousel-card-image_lax_1200_1200/..." }  
  ]  
}
```

VIDEO

```
json  
{  
  "adId": "771309874",  
  "format": "VIDEO",  
  "headline": "👉 Bekijk de Online Masterclass",  
  "videoUrl": "https://dms.licdn.com/playlist/vid/v2/.../progressive-servable-video/..."  
}
```

DOCUMENT

json

```
{  
  "adId": "685869434",  
  "format": "DOCUMENT",  
  "documentUrl": "https://media.licdn.com/dms/document/media/v2/.../document-sliced-pdf_multiton/...",  
  "imageUrls": [  
    "https://media.licdn.com/dms/image/v2/.../ads-document-images_1920/...",  
    "https://media.licdn.com/dms/image/v2/.../ads-document-images_1920/..."  
  ]  
}
```

MESSAGE (InMail — minimal data)

json

```
{  
  "adId": "943272674",  
  "format": "MESSAGE",  
  "body": "He %FIRSTNAME%,\n\nBenieuwd naar het uurtarief van consultancybureaus...",  
  "impressions": "5k-10k",  
  "targeting": { "language": "English", "location": "Netherlands" }  
}
```

SPOTLIGHT (Dynamic Ad)

json

```
{  
  "adId": "750972514",  
  "format": "SPOTLIGHT",  
  "body": "Simplicate brengt alle projectinformatie samen op één plek",  
  "headline": "👉 Het overzicht over projecten en omzet kwijt? 🔥 ",  
  "clickUrl": "https://www.simplicate.com/nl/ingenieurs/...",  
  "ctas": ["Plan een demo"]  
}
```

FOLLOW_COMPANY (Dynamic Ad — note different paidBy)

json

```
{
  "adId": "819500756",
  "format": "FOLLOW_COMPANY",
  "paidBy": "Ideuzo Com",
  "impressions": "100k-150k",
  "clickUrl": "https://www.linkedin.com/company/2027242/life",
  "ctas": ["Visit life"]
}
```

EVENT

json

```
{
  "adId": "1091001534",
  "format": "EVENT",
  "body": "Consultancy bureaus: Meer rust, meer structuur en meer grip op sales...",
  "clickUrl": "https://www.linkedin.com/events/winterofsalesmetrickaarsman.../",
  "ctas": ["View event"],
  "impressionsPerCountry": [
    { "country": "Netherlands", "impressions": "100%" },
    { "country": "Portugal", "impressions": "<1%" }
  ]
}
```

10. Key Principles

- Isolation:** Only `lib/apify/transform.ts` knows Apify field names. Everything else uses Prisma types.
- Upsert, don't duplicate:** Always match on `externalId` (LinkedIn's ad ID). Re-scraping updates existing records.
- Store raw where useful:** `mediaData` (JSON) preserves format-specific detail. `impressionsPerCountry` (JSON) keeps geo data.
- Budget safety:** Never run a scrape without checking monthly spend first.
- Track everything:** Every scrape creates a `ScrapeRun` record for debugging and cost tracking.